




—■— z context encoder    —●— h target encoder    —▲— p predictor



each point = one head, trained and tested on that condition. y axis truncated at 79%.